

# Comparison of GAN Architectures

Ryan Missel  
Jeffery Russell  
Kyle Rivenburgh





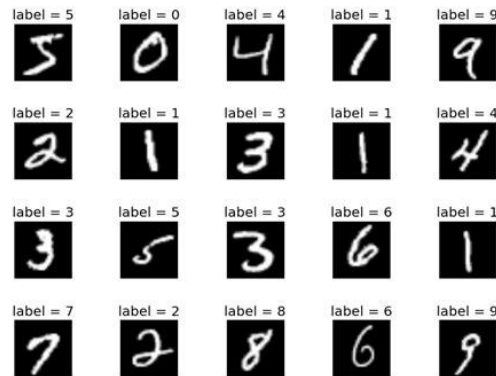
# Agenda

- Problem Statement
- What is a GAN?
- GAN Architectures
- Experiments/Results
- Conclusion



# Research Questions

- Which common GAN architecture performed the best on the MNIST Dataset?
- What are the quantitative differences in terms of:
  - Stability?
  - Result Quality?
- How does training differ between architectures?




**What is a GAN?**



# What is a GAN?

- Type of Neural Network called a Generative Adversarial Network
- Proposed by Ian Goodfellow in his 2014 thesis
- Consists of training 2 competing networks
  - One that learns to generate fake samples
  - One that learns to detect real vs. fake ones



# What is a GAN?

- Popular uses of GANs include:
  - Image Generation
  - Music Generation
  - Style Transfer
  - Video Prediction
  - Super Resolution
  - Text-to-image

StyleGAN 2 Results



Example of Style Transfer







# GAN Architectures



# GAN Architectures

- There are many different GAN architectures proposed in the field
- In this work, we compared 3 common ones:
  - Vanilla GAN
  - DCGAN
  - WGAN
- The differences primarily lie in layer shapes and design

Other architectures  
not explored








# GAN Architectures - Vanilla

- The vanilla architecture is just a dense network for both the model
- Generator:
  - 5 fully-connected layers
  - 128 -> 128 -> 256 -> 512 -> 1024
  - Each hidden layer is activated with a Leaky ReLU
  - The output layer is activated with a TanH for pixel scaling
- Discriminator:
  - 3 fully-connected layers
  - 1024 -> 512 -> 256 -> 1
  - Leaky ReLU for hidden layers, Sigmoid for output

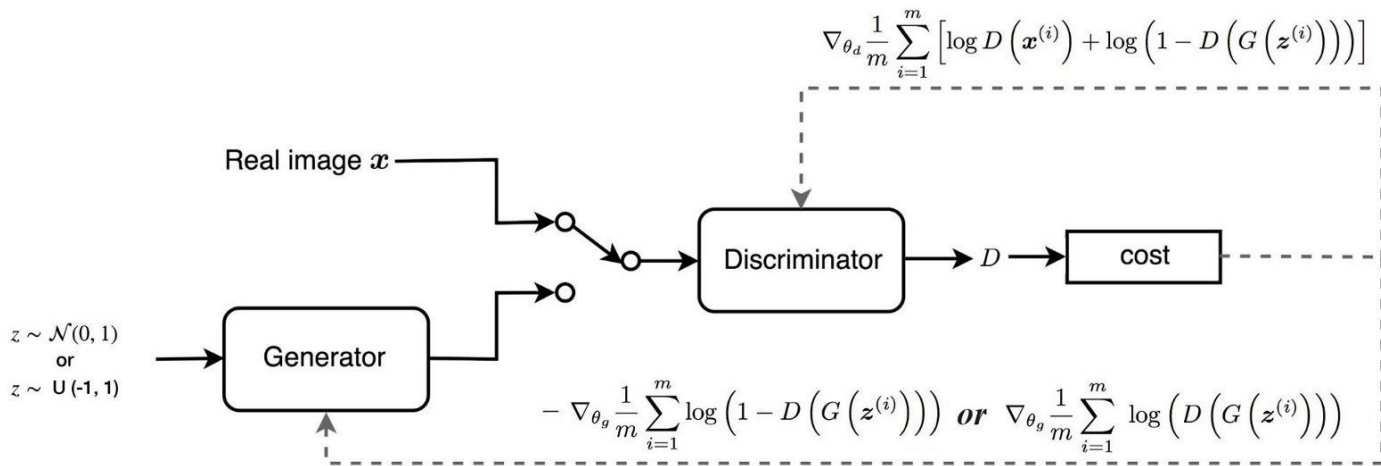
# GAN Architectures - DCGAN

- Stands for “Deep Convolutional GAN”
- Relates to having a mirrored deep convolutional network in both models
- Instead of fully-connected layers, convolution layers make sense for image work



# GAN Architectures - WGAN

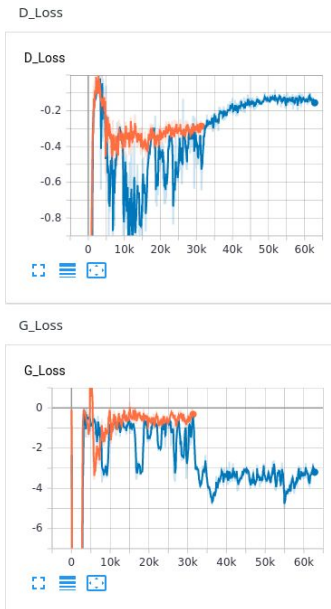
- Stands for “Wasserstein GAN”
- Changes the objective function of the discriminator to work off of “belief”
  - I.e. “An example is 60% real and 40% fake”
- Idea is to allow the Generator to minimize the distance between real and fake easier through a continuous output on the Discriminator side





# Experiments/Results

# Implementation




- Using identical hyperparameters and training methodologies, the three GAN architects were implemented in PyTorch.
- Tensorboard was used to view the performance measures of the discriminator and generator during training.
- We used the PyTorch built in loader for the MNIST dataset.




# Training Time


- Rates vastly differed between models, both for convergence and time per epoch
- All models were trained on a GTX2060
- Vanilla GAN
  - Quickest time per epoch
- WGAN
  - Similar amount of time to converge as the Vanilla GAN
  - Similar time per epoch
  - Makes sense since architecture similar, just loss function differences
- DCGAN
  - Slowest training time due to massive architecture

Results with  
no training






(a) 400 Batches




(b) 6000 Batches




(c) 187200 Batches

FIG. 3: GAN Results Sampled at Different Epochs


- DCGANs were able to converge on legible results with the fewest numbers of epochs.
- Our vanilla GAN took the longest to converge on legible outputs



(a) 400 Batches




(b) 6000 Batches



(c) 187200 Batches

FIG. 4: WGAN Results Sampled at Different Epochs

# Training Rate



(a) 400 Batches



(b) 6000 Batches



(c) 187200 Batches

FIG. 4: DCGAN Results Sampled at Different Epochs

# Perceived Quality

- To test the quality of the outputs, we sampled ratings from random participants
- Blind samples from a mix of the models
  - We took the last 6 results generated from each architecture -- after 200 epochs
  - Noted that the DCGAN had the best quality

GAN



WGAN



DCGAN




Vanilla GAN	WGAN	DCGAN
7	7	10
6	8	10
7	7	10
6	7	10
5	7	10
4	8	10



# Training Data Used


- Restricted the training data to 1/6th the size to test data needed
- Had a massive effect on quality of the samples
  - GAN/WGAN affected most heavily
  - DCGAN performed decently



(a) GAN



(b) WGAN



(c) DCGAN

FIG. 6: Results with one Sixth of Training Set and trained for 200 Epochs

# Questions?

Email us at:

Ryan Missel ([rxm7244@rit.edu](mailto:rxm7244@rit.edu))

Jeffery Russell ([jeffery@jrtechs.net](mailto:jeffery@jrtechs.net))

Kyle Rivenburgh ([ktr5669@rit.edu](mailto:ktr5669@rit.edu))

Code:

<https://github.com/jrtechs/CSCI-431-final-GANs>

